

Introduction to batch computing on the Flux cluster: Lmod modules

07 Oct 2016

1. Software modules

- Software modules enable you to control which software packages are available to use to prevent conflicts and control versions. We use a module program called Lmod, from the TACC
- How to find software that is installed

```
$ module keyword statistics
```

- To get some more information about a package

```
$ module spider stata
```

Note how it displays some information and gives you pointers to near misses. Not a perfect system, though

```
$ module spider R
```

that worked pretty well, but what about

```
$ module spider openmpi
```

scroll down until you find `openmpi/1.10.2/gcc` and there you will see the versions. Lmod considers `openmpi/1.10.2/gcc` to be the package name and the `4.8.5`, `4.9.3`, and `5.3.0` the version numbers.

- To make software available, you *load* it.

```
$ module load R
```

to see what is loaded, use

```
$ module list
```

to unload a loaded module, use

```
$ module unload R
```

to clear out all the loaded modules, use

```
$ module purge
```

- `module` commands need to be run once after logging in and before submitting a job to the queue. That insures that the settings are communicated to all the processes on all the nodes (using the `-v PBS`

flag).

- If you find you need the same module(s) all the time, you can create module *sets*. These are exactly what they sound like: A set of modules that can be treated as a unit. I recommend that you create sets after a fresh login and after purging loaded modules.

```
$ module purge
$ module load R
$ module load python-anaconda2
$ module list
```

Now you have a couple of modules loaded, save this as `stats` using

```
$ module save stats
```

To load a saved set, use

```
$ module restore set_name
```

To see the names of the save sets, use

```
$ module savelist
```

and to see what is in a set use

```
$ module describe set_name
```

Finally, there is a special set called `default` that will load automatically when you login. You create it the same way as any other set, but either name it `default` explicitly, or leave off the name.

Sets are saved as files in `~/.lmod.d`, and if you want to delete one, use the `rm` command.

- Make another set called `build-intel` that contains

```
intel/16.0.3 openmpi/1.6.5/intel/16.0.3
```

Now, purge your modules, restore the `stats` set. List your modules. Now restore your `build-intel` set. What are your loaded modules? So, restoring is not the same as loading a list; it's a list of all the modules that will be loaded after that state is restored.

- Modules can have dependencies

```
$ module spider gromacs
```

We want `gromacs/5.1.2/openmpi/1.10.2/gcc/4.9.3`. What does that name tell you? That's parsed as software/version pairs, where each package depends on the software/version(s) to its right. So, `gromacs` depends on `openmpi` which depends on `gcc`. Try loading that

```
$ module load gromacs/5.1.2/openmpi/1.10.2/gcc/4.9.3
```

Lo and behold, there are some other prerequisites, too. You can copy and paste

the dependencies onto a module load, and that's a set worth saving, by gum!

- Modules can have conflicts

```
$ module purge
$ module restore stats
$ module load python-epd
```

2. Now we have the pieces to run software that is already runnable, e.g., things that are installed on Flux and don't require anything additional to run, like R, Stata, Matlab.

- Example of developing a working script. Suppose we are embarking on a study of plant phylogeny. We have some data already stored in an R library. As the study progresses, the same statistics will need to be run as each round of data is added. So, first we determine, using a small sample, on the login node which commands are needed and to check that we have syntax correct.

```
$ module load R/3.1.1
$ R
```

```
> library(datasets)
> data(iris)
> summary(iris)
```

- Now that we have that working, we need to put that into a script and check that the batch version works. Enter those commands into `iris.R`

```
$ mkdir R-examples
$ cd R-examples
$ R CMD BATCH --no-restore --no-save iris.R
```

- Next, we need to make a PBS script to run the batch job. Copy your `template.pbs` to `iris.pbs`, fill in the necessary options to run the R script and submit it.

Optional section on compiling software

1. Compiling software

- Serial or shared memory software – one physical machine
- Use the compiler directly

```
$ cp -r /scratch/data/workshops/hpc101/code-examples .
$ gcc -o hello hello.c
$ ./hello
```

- To change compiler

```
$ module load intel
$ gcc -o hello hello.c
$ ./hello
```

- For this simple example, `gcc` and `intel-comp` can coexist, however it is good practice to unload the unused compiler if you can.
- Software you compile yourself is run the same way as commercial software
- Parallel or MPI-based software – more than one physical machine
- Uses compiler scripts instead of using the compiler directly
- Compiler and MPI version must match
 - Check with `module av openmpi`, which gives yet another view of the choice of modules; which modules contain `openmpi` somewhere in the name.

```
$ mpicc -o hello mpi-hello.c
$ mpirun -np 4 hello
```

- You must load the correct (matching) `openmpi` to run a program compiled with it
- When run from inside a batch job, you do not need `-np N` unless you use fewer processors than requested by the PBS script.
- To change compilers, you have to change both the compiler and the `openmpi`

```
$ module purge
$ module load gcc/4.9.3 openmpi/1.10.2/gcc/4.9.3
$ mpicc -o hello mpi-hello.c
$ mpirun -np 4 hello
```